

# HW2\_task1

March 15, 2023

This is a implementation of a neural network using PyTorch for a classification task on a generated dataset. The dataset consists of 1000 data points with two features, divided into four clusters. The network has one hidden layer consists of 10 neurons. The input is passed through a ReLU activation function before being passed to the output layer, which has a log-softmax activation. The network is trained using stochastic gradient descent (SGD) with a learning rate of 0.01 and a negative log-likelihood loss function (NLLLoss).

The training is done for 100 epochs, with the training loss and test loss being recorded at each epoch. The training and test losses are plotted over the number of epochs. The final accuracy of the trained model is evaluated on both the training and test sets.


The model achieves a high accuracy on both the training set (99%) and the test set (100%), indicating that it is able to generalize well to unseen data. The plot of the training and test losses shows that the model is not overfitting, as the test loss does not increase while the training loss decreases.

```
[ ]: import numpy as np
      from sklearn.datasets import make_blobs
      import matplotlib.pyplot as plt
      import torch.optim as optim
      import torch
      import torch.nn as nn

      # Generate 4 clusters of data points
      X, y = make_blobs(n_samples=1000, centers=4, n_features=2, random_state=42)

      # Randomly permute the dataset
      perm = np.random.permutation(X.shape[0])
      X = X[perm]
      y = y[perm]

      # Split the data into training and testing sets
      X_train, y_train = X[:900], y[:900]
      X_test, y_test = X[900:], y[900:]
      plt.scatter(X_train[:, 0], X_train[:, 1], c=y_train)
      plt.show()
```



```
[ ]: class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(2, 10)
        self.fc2 = nn.Linear(10, 4)

    def forward(self, x):
        x = torch.relu(self.fc1(x))
        x = self.fc2(x)
        return torch.log_softmax(x, dim=1)

net = Net()
```

```
[ ]: train_losses = []
    test_losses = []

    criterion = nn.NLLLoss()
    optimizer = optim.SGD(net.parameters(), lr=0.01)

    for epoch in range(100):
        running_loss = 0.0
        for i in range(0, X_train.shape[0], 10):
            inputs = torch.tensor(X_train[i:i+10], dtype=torch.float)
            labels = torch.tensor(y_train[i:i+10], dtype=torch.long)
```

```

optimizer.zero_grad()

outputs = net(inputs)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()

running_loss += loss.item()

# Record the average loss per batch for the training set
train_losses.append(running_loss / (X_train.shape[0] / 10))

# Compute the loss on the test set
with torch.no_grad():
    inputs = torch.tensor(X_test, dtype=torch.float)
    labels = torch.tensor(y_test, dtype=torch.long)
    outputs = net(inputs)
    test_loss = criterion(outputs, labels)
    test_losses.append(test_loss.item())

print('%d train loss: %.3f, test loss: %.3f' % (epoch + 1,
→train_losses[-1], test_losses[-1]))

```

```


[1] train loss: 0.700, test loss: 0.234
[2] train loss: 0.145, test loss: 0.094
[3] train loss: 0.072, test loss: 0.058
[4] train loss: 0.048, test loss: 0.042
[5] train loss: 0.036, test loss: 0.033
[6] train loss: 0.028, test loss: 0.027
[7] train loss: 0.024, test loss: 0.023
[8] train loss: 0.020, test loss: 0.020
[9] train loss: 0.018, test loss: 0.018
[10] train loss: 0.016, test loss: 0.016
[11] train loss: 0.014, test loss: 0.015
[12] train loss: 0.013, test loss: 0.013
[13] train loss: 0.012, test loss: 0.012
[14] train loss: 0.011, test loss: 0.012
[15] train loss: 0.011, test loss: 0.011
[16] train loss: 0.010, test loss: 0.010
[17] train loss: 0.009, test loss: 0.010
[18] train loss: 0.009, test loss: 0.009
[19] train loss: 0.008, test loss: 0.009
[20] train loss: 0.008, test loss: 0.008
[21] train loss: 0.008, test loss: 0.008
[22] train loss: 0.007, test loss: 0.008
[23] train loss: 0.007, test loss: 0.007
[24] train loss: 0.007, test loss: 0.007
[25] train loss: 0.007, test loss: 0.007

```

[26] train loss: 0.006, test loss: 0.007  
[27] train loss: 0.006, test loss: 0.006  
[28] train loss: 0.006, test loss: 0.006  
[29] train loss: 0.006, test loss: 0.006  
[30] train loss: 0.006, test loss: 0.006  
[31] train loss: 0.006, test loss: 0.006  
[32] train loss: 0.005, test loss: 0.006  
[33] train loss: 0.005, test loss: 0.005  
[34] train loss: 0.005, test loss: 0.005  
[35] train loss: 0.005, test loss: 0.005  
[36] train loss: 0.005, test loss: 0.005  
[37] train loss: 0.005, test loss: 0.005  
[38] train loss: 0.005, test loss: 0.005  
[39] train loss: 0.005, test loss: 0.005  
[40] train loss: 0.005, test loss: 0.005  
[41] train loss: 0.005, test loss: 0.005  
[42] train loss: 0.005, test loss: 0.004  
[43] train loss: 0.004, test loss: 0.004  
[44] train loss: 0.004, test loss: 0.004  
[45] train loss: 0.004, test loss: 0.004  
[46] train loss: 0.004, test loss: 0.004  
[47] train loss: 0.004, test loss: 0.004  
[48] train loss: 0.004, test loss: 0.004  
[49] train loss: 0.004, test loss: 0.004  
[50] train loss: 0.004, test loss: 0.004  
[51] train loss: 0.004, test loss: 0.004  
[52] train loss: 0.004, test loss: 0.004  
[53] train loss: 0.004, test loss: 0.004  
[54] train loss: 0.004, test loss: 0.004  
[55] train loss: 0.004, test loss: 0.004  
[56] train loss: 0.004, test loss: 0.004  
[57] train loss: 0.004, test loss: 0.003  
[58] train loss: 0.004, test loss: 0.003  
[59] train loss: 0.004, test loss: 0.003  
[60] train loss: 0.004, test loss: 0.003  
[61] train loss: 0.004, test loss: 0.003  
[62] train loss: 0.004, test loss: 0.003  
[63] train loss: 0.003, test loss: 0.003  
[64] train loss: 0.003, test loss: 0.003  
[65] train loss: 0.003, test loss: 0.003  
[66] train loss: 0.003, test loss: 0.003  
[67] train loss: 0.003, test loss: 0.003  
[68] train loss: 0.003, test loss: 0.003  
[69] train loss: 0.003, test loss: 0.003  
[70] train loss: 0.003, test loss: 0.003  
[71] train loss: 0.003, test loss: 0.003  
[72] train loss: 0.003, test loss: 0.003  
[73] train loss: 0.003, test loss: 0.003

```
[74] train loss: 0.003, test loss: 0.003
[75] train loss: 0.003, test loss: 0.003
[76] train loss: 0.003, test loss: 0.003
[77] train loss: 0.003, test loss: 0.003
[78] train loss: 0.003, test loss: 0.003
[79] train loss: 0.003, test loss: 0.003
[80] train loss: 0.003, test loss: 0.003
[81] train loss: 0.003, test loss: 0.003
[82] train loss: 0.003, test loss: 0.003
[83] train loss: 0.003, test loss: 0.003
[84] train loss: 0.003, test loss: 0.003
[85] train loss: 0.003, test loss: 0.003
[86] train loss: 0.003, test loss: 0.003
[87] train loss: 0.003, test loss: 0.003
[88] train loss: 0.003, test loss: 0.003
[89] train loss: 0.003, test loss: 0.003
[90] train loss: 0.003, test loss: 0.003
[91] train loss: 0.003, test loss: 0.003
[92] train loss: 0.003, test loss: 0.002
[93] train loss: 0.003, test loss: 0.002
[94] train loss: 0.003, test loss: 0.002
[95] train loss: 0.003, test loss: 0.002
[96] train loss: 0.003, test loss: 0.002
[97] train loss: 0.003, test loss: 0.002
[98] train loss: 0.003, test loss: 0.002
[99] train loss: 0.003, test loss: 0.002
[100] train loss: 0.003, test loss: 0.002
```

```
[ ]: plt.plot(train_losses, label='Train loss')
plt.plot(test_losses, label='Test loss')
plt.legend()
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.show()
```



```
[ ]: with torch.no_grad():
    inputs = torch.tensor(X_train, dtype=torch.float)
    labels = torch.tensor(y_train, dtype=torch.long)
    outputs = net(inputs)
    _, predicted = torch.max(outputs.data, 1)
    total = labels.size(0)
    correct = (predicted == labels).sum().item()

    print('Accuracy on the training set: %d %%' % (100 * correct / total))

    with torch.no_grad():
        inputs = torch.tensor(X_test, dtype=torch.float)
        labels = torch.tensor(y_test, dtype=torch.long)
        outputs = net(inputs)
        _, predicted = torch.max(outputs.data, 1)
        total = labels.size(0)
        correct = (predicted == labels).sum().item()

    print('Accuracy on the test set: %d %%' % (100 * correct / total))
```

Accuracy on the training set: 99 %  
 Accuracy on the test set: 100 %

## HW2\_task2

March 15, 2023

The code defines and trains a neural network model to learn the function  $f(x, y) = x^2 + xy + y^2$  on a random set of 5000 points generated from a uniform distribution on the plane  $[-10, 10] \times [-10, 10]$ . The data is split into training and testing sets with a 90/10 split. The neural network model has one hidden layer, which consists of 64 neurons, and ReLU activation. The model is trained using the mean squared error (MSE) loss function and the Adam optimizer with learning rate of 0.01. The model is trained for 200 epochs, and the batch size is set to 32. At each epoch, the model is trained on the training set and evaluated on the testing set. The training and testing losses are recorded and plotted as a function of the training time. The final training and testing losses are also reported at the end of the training.

```
[ ]: import numpy as np
from sklearn.model_selection import train_test_split
import torch
import torch.nn as nn
import torch.optim as optim
import matplotlib.pyplot as plt

# Generate 5000 random points (x, y) from a uniform distribution on [-10, 10] x
↳ [-10, 10]
x = np.random.uniform(low=-10, high=10, size=(5000, 2))

# Define the function f(x, y) = x^2 + xy + y^2
y = x[:, 0]**2 + x[:, 0]*x[:, 1] + x[:, 1]**2

# Reshape y to a column vector
y = y.reshape(-1, 1)

# Split the data into training and testing sets with a 90/10 split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.1,
↳ random_state=42)

# Move data to GPU
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
x_train = torch.tensor(x_train, dtype=torch.float32).to(device)
x_test = torch.tensor(x_test, dtype=torch.float32).to(device)
y_train = torch.tensor(y_train, dtype=torch.float32).to(device)
y_test = torch.tensor(y_test, dtype=torch.float32).to(device)
```

```

[ ]: # Define the model
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(2, 64)
        self.fc2 = nn.Linear(64, 1)

    def forward(self, x):
        x = torch.relu(self.fc1(x))
        x = self.fc2(x)
        return x

model = Net().to(device)

# Define the loss function and optimizer
criterion = nn.MSELoss()
optimizer = optim.Adam(model.parameters(), lr=0.001)

# Train the model on GPU
train_loss = []
test_loss = []
for epoch in range(200):
    running_train_loss = 0.0
    running_test_loss = 0.0
    for i in range(0, len(x_train), 32):
        # Train
        inputs = x_train[i:i+32]
        labels = y_train[i:i+32]

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_train_loss += loss.item() * inputs.size(0)

    # Evaluate on test set
    with torch.no_grad():
        for i in range(0, len(x_test), 32):
            inputs = x_test[i:i+32]
            labels = y_test[i:i+32]

            outputs = model(inputs)
            loss = criterion(outputs, labels)

            running_test_loss += loss.item() * inputs.size(0)

```



```

epoch_train_loss = running_train_loss / len(x_train)
epoch_test_loss = running_test_loss / len(x_test)
train_loss.append(epoch_train_loss)
test_loss.append(epoch_test_loss)

print(f"Epoch {epoch+1:3d} | Train loss: {epoch_train_loss:.6f} | Test loss:␣
→{epoch_test_loss:.6f}")

```

```

Epoch 1 | Train loss: 5856.777300 | Test loss: 4491.517270
Epoch 2 | Train loss: 3336.986687 | Test loss: 1976.690430
Epoch 3 | Train loss: 1346.534081 | Test loss: 773.401029
Epoch 4 | Train loss: 654.550887 | Test loss: 535.588102
Epoch 5 | Train loss: 520.628203 | Test loss: 484.840931
Epoch 6 | Train loss: 480.223928 | Test loss: 455.964436
Epoch 7 | Train loss: 454.086652 | Test loss: 432.844595
Epoch 8 | Train loss: 432.282651 | Test loss: 412.202475
Epoch 9 | Train loss: 412.338537 | Test loss: 392.910503
Epoch 10 | Train loss: 393.490713 | Test loss: 374.663848
Epoch 11 | Train loss: 375.502190 | Test loss: 357.289534
Epoch 12 | Train loss: 358.242575 | Test loss: 340.659789
Epoch 13 | Train loss: 341.623235 | Test loss: 324.652339
Epoch 14 | Train loss: 325.569345 | Test loss: 309.188141
Epoch 15 | Train loss: 310.033492 | Test loss: 294.226150
Epoch 16 | Train loss: 294.933388 | Test loss: 279.632452
Epoch 17 | Train loss: 280.207497 | Test loss: 265.407010
Epoch 18 | Train loss: 265.828967 | Test loss: 251.509280
Epoch 19 | Train loss: 251.781757 | Test loss: 237.894889
Epoch 20 | Train loss: 238.014035 | Test loss: 224.537107
Epoch 21 | Train loss: 224.479556 | Test loss: 211.394884
Epoch 22 | Train loss: 211.163343 | Test loss: 198.448694
Epoch 23 | Train loss: 198.106658 | Test loss: 185.799515
Epoch 24 | Train loss: 185.425612 | Test loss: 173.663835
Epoch 25 | Train loss: 173.280431 | Test loss: 162.040337
Epoch 26 | Train loss: 161.688923 | Test loss: 150.940518
Epoch 27 | Train loss: 150.640622 | Test loss: 140.399422
Epoch 28 | Train loss: 140.174912 | Test loss: 130.459874
Epoch 29 | Train loss: 130.331551 | Test loss: 121.172540
Epoch 30 | Train loss: 121.128282 | Test loss: 112.535347
Epoch 31 | Train loss: 112.557181 | Test loss: 104.525226
Epoch 32 | Train loss: 104.616572 | Test loss: 97.150884
Epoch 33 | Train loss: 97.302488 | Test loss: 90.414604
Epoch 34 | Train loss: 90.598669 | Test loss: 84.263615
Epoch 35 | Train loss: 84.472387 | Test loss: 78.695864
Epoch 36 | Train loss: 78.888396 | Test loss: 73.639963
Epoch 37 | Train loss: 73.805275 | Test loss: 69.054317
Epoch 38 | Train loss: 69.187067 | Test loss: 64.907319
Epoch 39 | Train loss: 64.979316 | Test loss: 61.147170

```

Epoch	40		Train loss:	61.143590		Test loss:	57.732182
Epoch	41		Train loss:	57.645259		Test loss:	54.617910
Epoch	42		Train loss:	54.442421		Test loss:	51.785044
Epoch	43		Train loss:	51.498891		Test loss:	49.167587
Epoch	44		Train loss:	48.774182		Test loss:	46.739027
Epoch	45		Train loss:	46.251618		Test loss:	44.471678
Epoch	46		Train loss:	43.902727		Test loss:	42.352100
Epoch	47		Train loss:	41.697260		Test loss:	40.342971
Epoch	48		Train loss:	39.613388		Test loss:	38.435758
Epoch	49		Train loss:	37.637631		Test loss:	36.609374
Epoch	50		Train loss:	35.763128		Test loss:	34.863918
Epoch	51		Train loss:	33.982912		Test loss:	33.194711
Epoch	52		Train loss:	32.276577		Test loss:	31.602441
Epoch	53		Train loss:	30.638842		Test loss:	30.088493
Epoch	54		Train loss:	29.069500		Test loss:	28.639283
Epoch	55		Train loss:	27.570091		Test loss:	27.252241
Epoch	56		Train loss:	26.135863		Test loss:	25.909296
Epoch	57		Train loss:	24.754468		Test loss:	24.612653
Epoch	58		Train loss:	23.423801		Test loss:	23.351341
Epoch	59		Train loss:	22.144631		Test loss:	22.134856
Epoch	60		Train loss:	20.914554		Test loss:	20.948720
Epoch	61		Train loss:	19.731542		Test loss:	19.812609
Epoch	62		Train loss:	18.595496		Test loss:	18.730342
Epoch	63		Train loss:	17.505996		Test loss:	17.689298
Epoch	64		Train loss:	16.460768		Test loss:	16.696817
Epoch	65		Train loss:	15.462396		Test loss:	15.752924
Epoch	66		Train loss:	14.510775		Test loss:	14.848936
Epoch	67		Train loss:	13.606289		Test loss:	13.991003
Epoch	68		Train loss:	12.751478		Test loss:	13.184846
Epoch	69		Train loss:	11.945278		Test loss:	12.417578
Epoch	70		Train loss:	11.185289		Test loss:	11.677351
Epoch	71		Train loss:	10.469312		Test loss:	10.974829
Epoch	72		Train loss:	9.794736		Test loss:	10.309331
Epoch	73		Train loss:	9.155422		Test loss:	9.677706
Epoch	74		Train loss:	8.553867		Test loss:	9.087158
Epoch	75		Train loss:	7.991909		Test loss:	8.531565
Epoch	76		Train loss:	7.461555		Test loss:	8.008160
Epoch	77		Train loss:	6.963404		Test loss:	7.518998
Epoch	78		Train loss:	6.497920		Test loss:	7.061846
Epoch	79		Train loss:	6.067349		Test loss:	6.635368
Epoch	80		Train loss:	5.668870		Test loss:	6.236466
Epoch	81		Train loss:	5.298138		Test loss:	5.866126
Epoch	82		Train loss:	4.952974		Test loss:	5.515210
Epoch	83		Train loss:	4.632550		Test loss:	5.190822
Epoch	84		Train loss:	4.333771		Test loss:	4.887172
Epoch	85		Train loss:	4.055543		Test loss:	4.603534
Epoch	86		Train loss:	3.795846		Test loss:	4.338200
Epoch	87		Train loss:	3.554163		Test loss:	4.092868


Epoch	88		Train loss:	3.329803		Test loss:	3.862264
Epoch	89		Train loss:	3.121820		Test loss:	3.646956
Epoch	90		Train loss:	2.928700		Test loss:	3.444850
Epoch	91		Train loss:	2.749588		Test loss:	3.259068
Epoch	92		Train loss:	2.583342		Test loss:	3.085827
Epoch	93		Train loss:	2.429469		Test loss:	2.925901
Epoch	94		Train loss:	2.286190		Test loss:	2.773680
Epoch	95		Train loss:	2.152559		Test loss:	2.629618
Epoch	96		Train loss:	2.027891		Test loss:	2.496135
Epoch	97		Train loss:	1.911489		Test loss:	2.369933
Epoch	98		Train loss:	1.802684		Test loss:	2.252359
Epoch	99		Train loss:	1.702525		Test loss:	2.140007
Epoch	100		Train loss:	1.608868		Test loss:	2.034733
Epoch	101		Train loss:	1.521517		Test loss:	1.933120
Epoch	102		Train loss:	1.439643		Test loss:	1.834898
Epoch	103		Train loss:	1.363022		Test loss:	1.740068
Epoch	104		Train loss:	1.291177		Test loss:	1.647117
Epoch	105		Train loss:	1.223872		Test loss:	1.557752
Epoch	106		Train loss:	1.161032		Test loss:	1.472091
Epoch	107		Train loss:	1.101835		Test loss:	1.392568
Epoch	108		Train loss:	1.046317		Test loss:	1.319446
Epoch	109		Train loss:	0.994280		Test loss:	1.249531
Epoch	110		Train loss:	0.944967		Test loss:	1.182960
Epoch	111		Train loss:	0.898102		Test loss:	1.118356
Epoch	112		Train loss:	0.854062		Test loss:	1.058363
Epoch	113		Train loss:	0.812829		Test loss:	1.004405
Epoch	114		Train loss:	0.774280		Test loss:	0.952894
Epoch	115		Train loss:	0.738010		Test loss:	0.902732
Epoch	116		Train loss:	0.703632		Test loss:	0.857586
Epoch	117		Train loss:	0.670648		Test loss:	0.816165
Epoch	118		Train loss:	0.639734		Test loss:	0.777989
Epoch	119		Train loss:	0.610748		Test loss:	0.741079
Epoch	120		Train loss:	0.583565		Test loss:	0.708820
Epoch	121		Train loss:	0.558340		Test loss:	0.679649
Epoch	122		Train loss:	0.534889		Test loss:	0.652153
Epoch	123		Train loss:	0.512590		Test loss:	0.626799
Epoch	124		Train loss:	0.492039		Test loss:	0.602061
Epoch	125		Train loss:	0.472628		Test loss:	0.578857
Epoch	126		Train loss:	0.454283		Test loss:	0.555963
Epoch	127		Train loss:	0.436606		Test loss:	0.535173
Epoch	128		Train loss:	0.419995		Test loss:	0.516468
Epoch	129		Train loss:	0.404159		Test loss:	0.497006
Epoch	130		Train loss:	0.389545		Test loss:	0.480207
Epoch	131		Train loss:	0.375571		Test loss:	0.463534
Epoch	132		Train loss:	0.362134		Test loss:	0.447794
Epoch	133		Train loss:	0.349353		Test loss:	0.432452
Epoch	134		Train loss:	0.336874		Test loss:	0.417220
Epoch	135		Train loss:	0.325167		Test loss:	0.400819

Epoch 136		Train loss: 0.314104		Test loss: 0.386807
Epoch 137		Train loss: 0.303613		Test loss: 0.372552
Epoch 138		Train loss: 0.293398		Test loss: 0.362485
Epoch 139		Train loss: 0.283860		Test loss: 0.351804
Epoch 140		Train loss: 0.274696		Test loss: 0.342148
Epoch 141		Train loss: 0.266033		Test loss: 0.333181
Epoch 142		Train loss: 0.257792		Test loss: 0.322504
Epoch 143		Train loss: 0.250088		Test loss: 0.312897
Epoch 144		Train loss: 0.243173		Test loss: 0.306793
Epoch 145		Train loss: 0.236560		Test loss: 0.298758
Epoch 146		Train loss: 0.230513		Test loss: 0.290423
Epoch 147		Train loss: 0.224906		Test loss: 0.284053
Epoch 148		Train loss: 0.219685		Test loss: 0.276969
Epoch 149		Train loss: 0.214624		Test loss: 0.269828
Epoch 150		Train loss: 0.209836		Test loss: 0.263546
Epoch 151		Train loss: 0.205451		Test loss: 0.257618
Epoch 152		Train loss: 0.201117		Test loss: 0.251383
Epoch 153		Train loss: 0.196896		Test loss: 0.245665
Epoch 154		Train loss: 0.192964		Test loss: 0.240968
Epoch 155		Train loss: 0.189074		Test loss: 0.236763
Epoch 156		Train loss: 0.185386		Test loss: 0.232548
Epoch 157		Train loss: 0.181691		Test loss: 0.227611
Epoch 158		Train loss: 0.178445		Test loss: 0.223676
Epoch 159		Train loss: 0.175343		Test loss: 0.219192
Epoch 160		Train loss: 0.172319		Test loss: 0.214863
Epoch 161		Train loss: 0.169436		Test loss: 0.210761
Epoch 162		Train loss: 0.166712		Test loss: 0.206577
Epoch 163		Train loss: 0.164055		Test loss: 0.203197
Epoch 164		Train loss: 0.161630		Test loss: 0.198900
Epoch 165		Train loss: 0.159186		Test loss: 0.194594
Epoch 166		Train loss: 0.156533		Test loss: 0.190014
Epoch 167		Train loss: 0.153946		Test loss: 0.185499
Epoch 168		Train loss: 0.151695		Test loss: 0.181589
Epoch 169		Train loss: 0.149445		Test loss: 0.175769
Epoch 170		Train loss: 0.147176		Test loss: 0.171181
Epoch 171		Train loss: 0.144845		Test loss: 0.165846
Epoch 172		Train loss: 0.142833		Test loss: 0.162096
Epoch 173		Train loss: 0.140810		Test loss: 0.157704
Epoch 174		Train loss: 0.138775		Test loss: 0.153456
Epoch 175		Train loss: 0.136540		Test loss: 0.148475
Epoch 176		Train loss: 0.134247		Test loss: 0.143552
Epoch 177		Train loss: 0.132160		Test loss: 0.139526
Epoch 178		Train loss: 0.130086		Test loss: 0.135793
Epoch 179		Train loss: 0.128307		Test loss: 0.132898
Epoch 180		Train loss: 0.126734		Test loss: 0.129659
Epoch 181		Train loss: 0.125260		Test loss: 0.126910
Epoch 182		Train loss: 0.123871		Test loss: 0.124557
Epoch 183		Train loss: 0.122560		Test loss: 0.122412

Epoch 184		Train loss: 0.121229		Test loss: 0.120414
Epoch 185		Train loss: 0.120076		Test loss: 0.118468
Epoch 186		Train loss: 0.118763		Test loss: 0.115934
Epoch 187		Train loss: 0.117528		Test loss: 0.114579
Epoch 188		Train loss: 0.116572		Test loss: 0.113097
Epoch 189		Train loss: 0.115406		Test loss: 0.111992
Epoch 190		Train loss: 0.114441		Test loss: 0.111067
Epoch 191		Train loss: 0.113392		Test loss: 0.109863
Epoch 192		Train loss: 0.112567		Test loss: 0.108369
Epoch 193		Train loss: 0.111612		Test loss: 0.107399
Epoch 194		Train loss: 0.110757		Test loss: 0.106470
Epoch 195		Train loss: 0.109925		Test loss: 0.105691
Epoch 196		Train loss: 0.108963		Test loss: 0.104535
Epoch 197		Train loss: 0.108087		Test loss: 0.103260
Epoch 198		Train loss: 0.107124		Test loss: 0.102411
Epoch 199		Train loss: 0.106323		Test loss: 0.101363
Epoch 200		Train loss: 0.105617		Test loss: 0.100108

```
[ ]: # Plot the train and test loss as a function of the training time
plt.plot(train_loss, label='Train loss')
plt.plot(test_loss, label='Test loss')
plt.legend()
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.show()

# Report the final training and test loss
print(f"Final train loss: {train_loss[-1]:.6f}")
print(f"Final test loss: {test_loss[-1]:.6f}")
```



Final train loss: 0.105617

Final test loss: 0.100108

# HW2\_task3

March 15, 2023

## 1 Task 3

This code defines a Siamese Network in PyTorch. The Siamese network is a type of neural network architecture that is typically used for tasks involving image similarity estimation, such as face recognition, signature verification, and image retrieval. The Siamese network takes two input images and feeds them through two identical networks with shared weights, which produces two feature vectors. These feature vectors are concatenated and passed through a linear layer to produce a similarity score between the two input images.

This implementation of the Siamese network uses the ResNet-18 model as the feature extractor, and it is trained on only 10% of the MNIST training dataset, and use only 10% of MNIST test data set as the test data set. The code defines a PyTorch dataset class called APP\_MATCHER, which groups the MNIST dataset examples based on class. For every example, the **getitem** method selects two images. For positive examples, two images are selected from the same class and the label is set to 1, while for negative examples, two images are selected from different classes, and the label is set to 0. The forward method of the SiameseNetwork class takes two input images and returns the similarity score between them.

```
[1]: from __future__ import print_function
import argparse, random, copy
import numpy as np

import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
from torch.utils.data import Dataset
from torchvision import datasets
from torchvision import transforms as T
from torch.optim.lr_scheduler import StepLR
```

```
[2]: class SiameseNetwork(nn.Module):
    """
    Siamese network for image similarity estimation.
    The network is composed of two identical networks, one for each input.
    The output of each network is concatenated and passed to a linear layer.
    The output of the linear layer passed through a sigmoid function.
```

*"FaceNet" <<https://arxiv.org/pdf/1503.03832.pdf>>`\_ is a variant of the`  
→ Siamese network.*

*This implementation varies from FaceNet as we use the `ResNet-18` model`  
→ from*

*"Deep Residual Learning for Image Recognition" <<https://arxiv.org/pdf/1512.03385.pdf>>`\_ as our feature extractor.*

*In addition, we aren't using `TripletLoss` as the MNIST dataset is`  
→ simple, so `BCELoss` can do the trick.*

```
"""
def __init__(self):
    super(SiameseNetwork, self).__init__()
    # get resnet model
    self.resnet = torchvision.models.resnet18(weights=None)

    # over-write the first conv layer to be able to read MNIST images
    # as resnet18 reads (3,x,x) where 3 is RGB channels
    # whereas MNIST has (1,x,x) where 1 is a gray-scale channel
    self.resnet.conv1 = nn.Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2),
→padding=(3, 3), bias=False)
    self.fc_in_features = self.resnet.fc.in_features

    # remove the last layer of resnet18 (linear layer which is before`
→avgpool layer)
    self.resnet = torch.nn.Sequential(*(list(self.resnet.children())[:-1]))

    # add linear layers to compare between the features of the two images
    self.fc = nn.Sequential(
        nn.Linear(self.fc_in_features * 2, 256),
        nn.ReLU(inplace=True),
        nn.Linear(256, 1),
    )

    self.sigmoid = nn.Sigmoid()

    # initialize the weights
    self.resnet.apply(self.init_weights)
    self.fc.apply(self.init_weights)

def init_weights(self, m):
    if isinstance(m, nn.Linear):
        torch.nn.init.xavier_uniform_(m.weight)
        m.bias.data.fill_(0.01)

def forward_once(self, x):
    output = self.resnet(x)
    output = output.view(output.size()[0], -1)
```



```

    return output

def forward(self, input1, input2):
    # get two images' features
    output1 = self.forward_once(input1)
    output2 = self.forward_once(input2)

    # concatenate both images' features
    output = torch.cat((output1, output2), 1)

    # pass the concatenation to the linear layers
    output = self.fc(output)

    # pass the out of the linear layers to sigmoid layer
    output = self.sigmoid(output)

    return output

```

```

[3]: class APP_MATCHER(Dataset):
    def __init__(self, root, train, download=False):
        super(APP_MATCHER, self).__init__()

        # get MNIST dataset
        self.dataset = datasets.MNIST(root, train=train, download=download)

        # as `self.dataset.data`'s shape is (Nx28x28), where N is the number of
        # examples in MNIST dataset, a single example has the dimensions of
        # (28x28) for (WxH), where W and H are the width and the height of the
        ↪ image.
        # However, every example should have (CxWxH) dimensions where C is the
        ↪ number
        # of channels to be passed to the network. As MNIST contains gray-scale
        ↪ images,
        # we add an additional dimension to corresponds to the number of
        ↪ channels.
        self.data = self.dataset.data.unsqueeze(1).clone()

        self.group_examples()

    def group_examples(self):
        """
        To ease the accessibility of data based on the class, we will use
        ↪ `group_examples` to group
        examples based on class.

        Every key in `grouped_examples` corresponds to a class in MNIST
        ↪ dataset. For every key in

```

```

        `grouped_examples`, every value will conform to all of the indices
→for the MNIST
        dataset examples that correspond to that key.
        """

        # get the targets from MNIST dataset
        np_arr = np.array(self.dataset.targets.clone())

        # group examples based on class
        self.grouped_examples = {}
        for i in range(0,10):
            self.grouped_examples[i] = np.where((np_arr==i))[0]

    def __len__(self):
        return self.data.shape[0]

    def __getitem__(self, index):
        """
        For every example, we will select two images. There are two cases,
        positive and negative examples. For positive examples, we will have
→two
        images from the same class. For negative examples, we will have two
→images
        from different classes.
        Given an index, if the index is even, we will pick the second image
→from the same class,
        but it won't be the same image we chose for the first class. This is
→used to ensure the positive
        example isn't trivial as the network would easily distinguish the
→similarity between same images. However,
        if the network were given two different images from the same class,
→the network will need to learn
        the similarity between two different images representing the same
→class. If the index is odd, we will
        pick the second image from a different class than the first image.
        """

        # pick some random class for the first image
        selected_class = random.randint(0, 9)

        # pick a random index for the first image in the grouped indices based
→of the label
        # of the class
        random_index_1 = random.randint(0, self.grouped_examples[selected_class].
→shape[0]-1)

```

```

# pick the index to get the first image
index_1 = self.grouped_examples[selected_class][random_index_1]

# get the first image
image_1 = self.data[index_1].clone().float()

# same class
if index % 2 == 0:
    # pick a random index for the second image
    random_index_2 = random.randint(0, self.
→grouped_examples[selected_class].shape[0]-1)

    # ensure that the index of the second image isn't the same as the
→first image
    while random_index_2 == random_index_1:
        random_index_2 = random.randint(0, self.
→grouped_examples[selected_class].shape[0]-1)

    # pick the index to get the second image
    index_2 = self.grouped_examples[selected_class][random_index_2]

    # get the second image
    image_2 = self.data[index_2].clone().float()

    # set the label for this example to be positive (1)
    target = torch.tensor(1, dtype=torch.float)

# different class
else:
    # pick a random class
    other_selected_class = random.randint(0, 9)

    # ensure that the class of the second image isn't the same as the
→first image
    while other_selected_class == selected_class:
        other_selected_class = random.randint(0, 9)

    # pick a random index for the second image in the grouped indices
→based of the label
    # of the class
    random_index_2 = random.randint(0, self.
→grouped_examples[other_selected_class].shape[0]-1)

    # pick the index to get the second image
    index_2 = self.grouped_examples[other_selected_class][random_index_2]

```

```

        # get the second image
        image_2 = self.data[index_2].clone().float()

        # set the label for this example to be negative (0)
        target = torch.tensor(0, dtype=torch.float)

    return image_1, image_2, target

```

```

[ ]: def train(args, model, device, train_loader, optimizer, epoch):
    model.train()

    # we aren't using `TripletLoss` as the MNIST dataset is simple, so `BCELoss`
    ↪ can do the trick.
    criterion = nn.BCELoss()

    for batch_idx, (images_1, images_2, targets) in enumerate(train_loader):
        images_1, images_2, targets = images_1.to(device), images_2.to(device),
        ↪ targets.to(device)
        optimizer.zero_grad()
        outputs = model(images_1, images_2).squeeze()
        loss = criterion(outputs, targets)
        loss.backward()
        optimizer.step()
        if batch_idx % args.log_interval == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(images_1), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))
            if args.dry_run:
                break

def test(model, device, test_loader):
    model.eval()
    test_loss = 0
    correct = 0

    # we aren't using `TripletLoss` as the MNIST dataset is simple, so `BCELoss`
    ↪ can do the trick.
    criterion = nn.BCELoss()

    with torch.no_grad():
        for (images_1, images_2, targets) in test_loader:
            images_1, images_2, targets = images_1.to(device), images_2.
            ↪ to(device), targets.to(device)
            outputs = model(images_1, images_2).squeeze()

```

```

        test_loss += criterion(outputs, targets).sum().item() # sum up
↪batch loss
        pred = torch.where(outputs > 0.5, 1, 0) # get the index of the max
↪log-probability
        correct += pred.eq(targets.view_as(pred)).sum().item()

test_loss /= len(test_loader.dataset)

# for the 1st epoch, the average loss is 0.0001 and the accuracy 97-98%
# using default settings. After completing the 10th epoch, the average
# loss is 0.0000 and the accuracy 99.5-100% using default settings.
print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.0f}%)\n'.format(
    test_loss, correct, len(test_loader.dataset),
    100. * correct / len(test_loader.dataset)))

```

```

[5]: def main():
    # Training settings
    parser = argparse.ArgumentParser(description='PyTorch Siamese network
↪Example')
    parser.add_argument('--batch-size', type=int, default=64, metavar='N',
                        help='input batch size for training (default: 64)')
    parser.add_argument('--test-batch-size', type=int, default=1000, metavar='N',
                        help='input batch size for testing (default: 1000)')
    parser.add_argument('--epochs', type=int, default=14, metavar='N',
                        help='number of epochs to train (default: 14)')
    parser.add_argument('--lr', type=float, default=1.0, metavar='LR',
                        help='learning rate (default: 1.0)')
    parser.add_argument('--gamma', type=float, default=0.7, metavar='M',
                        help='Learning rate step gamma (default: 0.7)')
    parser.add_argument('--no-cuda', action='store_true', default=False,
                        help='disables CUDA training')
    parser.add_argument('--no-mps', action='store_true', default=False,
                        help='disables macOS GPU training')
    parser.add_argument('--dry-run', action='store_true', default=False,
                        help='quickly check a single pass')
    parser.add_argument('--seed', type=int, default=1, metavar='S',
                        help='random seed (default: 1)')
    parser.add_argument('--log-interval', type=int, default=10, metavar='N',
                        help='how many batches to wait before logging training
↪status')
    parser.add_argument('--save-model', action='store_true', default=False,
                        help='For Saving the current Model')
    parser.add_argument('-f')
    args = parser.parse_args()

    use_cuda = not args.no_cuda and torch.cuda.is_available()
    use_mps = not args.no_mps and torch.backends.mps.is_available()

```

```

torch.manual_seed(args.seed)

if use_cuda:
    device = torch.device("cuda")
elif use_mps:
    device = torch.device("mps")
else:
    device = torch.device("cpu")

train_kwargs = {'batch_size': args.batch_size}
test_kwargs = {'batch_size': args.test_batch_size}
if use_cuda:
    cuda_kwargs = {'num_workers': 1,
                   'pin_memory': True,
                   'shuffle': True}
    train_kwargs.update(cuda_kwargs)
    test_kwargs.update(cuda_kwargs)

train_dataset = APP_MATCHER('../data', train=True, download=True)
train_indices = torch.randperm(len(train_dataset))[:len(train_dataset)//10]
train_dataset = torch.utils.data.Subset(train_dataset, train_indices)
train_loader = torch.utils.data.DataLoader(train_dataset, **train_kwargs)

test_dataset = APP_MATCHER('../data', train=False, download=True)
test_indices = torch.randperm(len(test_dataset))[:len(test_dataset)//10]
test_dataset = torch.utils.data.Subset(test_dataset, test_indices)
test_loader = torch.utils.data.DataLoader(test_dataset, **test_kwargs)

model = SiameseNetwork().to(device)
optimizer = optim.Adadelta(model.parameters(), lr=args.lr)

scheduler = StepLR(optimizer, step_size=1, gamma=args.gamma)
for epoch in range(1, 10):
    train(args, model, device, train_loader, optimizer, epoch)
    test(model, device, test_loader)
    scheduler.step()

if args.save_model:
    torch.save(model.state_dict(), "siamese_network.pt")

if __name__ == '__main__':
    main()

```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>  
 Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to  
 ../data/MNIST/raw/train-images-idx3-ubyte.gz

```

0%|          | 0/9912422 [00:00<?, ?it/s]
Extracting ../data/MNIST/raw/train-images-idx3-ubyte.gz to ../data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to
../data/MNIST/raw/train-labels-idx1-ubyte.gz

0%|          | 0/28881 [00:00<?, ?it/s]
Extracting ../data/MNIST/raw/train-labels-idx1-ubyte.gz to ../data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to
../data/MNIST/raw/t10k-images-idx3-ubyte.gz

0%|          | 0/1648877 [00:00<?, ?it/s]
Extracting ../data/MNIST/raw/t10k-images-idx3-ubyte.gz to ../data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to
../data/MNIST/raw/t10k-labels-idx1-ubyte.gz

0%|          | 0/4542 [00:00<?, ?it/s]
Extracting ../data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ../data/MNIST/raw

Train Epoch: 1 [0/6000 (0%)]    Loss: 1.175190
Train Epoch: 1 [640/6000 (11%)] Loss: 0.682234
Train Epoch: 1 [1280/6000 (21%)]    Loss: 0.652923
Train Epoch: 1 [1920/6000 (32%)]    Loss: 0.623879
Train Epoch: 1 [2560/6000 (43%)]    Loss: 0.627442
Train Epoch: 1 [3200/6000 (53%)]    Loss: 0.501702
Train Epoch: 1 [3840/6000 (64%)]    Loss: 0.323163
Train Epoch: 1 [4480/6000 (74%)]    Loss: 0.534408
Train Epoch: 1 [5120/6000 (85%)]    Loss: 0.411011
Train Epoch: 1 [5760/6000 (96%)]    Loss: 0.383505

Test set: Average loss: 0.0005, Accuracy: 801/1000 (80%)

Train Epoch: 2 [0/6000 (0%)]    Loss: 0.448275
Train Epoch: 2 [640/6000 (11%)] Loss: 0.414833
Train Epoch: 2 [1280/6000 (21%)]    Loss: 0.232005
Train Epoch: 2 [1920/6000 (32%)]    Loss: 0.364514
Train Epoch: 2 [2560/6000 (43%)]    Loss: 0.338736
Train Epoch: 2 [3200/6000 (53%)]    Loss: 0.262378
Train Epoch: 2 [3840/6000 (64%)]    Loss: 0.345625
Train Epoch: 2 [4480/6000 (74%)]    Loss: 0.234799
Train Epoch: 2 [5120/6000 (85%)]    Loss: 0.162417
Train Epoch: 2 [5760/6000 (96%)]    Loss: 0.205213

```

Test set: Average loss: 0.0003, Accuracy: 906/1000 (91%)

Train Epoch: 3	[0/6000 (0%)]	Loss: 0.191834
Train Epoch: 3	[640/6000 (11%)]	Loss: 0.150406
Train Epoch: 3	[1280/6000 (21%)]	Loss: 0.250022
Train Epoch: 3	[1920/6000 (32%)]	Loss: 0.151487
Train Epoch: 3	[2560/6000 (43%)]	Loss: 0.104292
Train Epoch: 3	[3200/6000 (53%)]	Loss: 0.106918
Train Epoch: 3	[3840/6000 (64%)]	Loss: 0.143443
Train Epoch: 3	[4480/6000 (74%)]	Loss: 0.064529
Train Epoch: 3	[5120/6000 (85%)]	Loss: 0.227299
Train Epoch: 3	[5760/6000 (96%)]	Loss: 0.222159

Test set: Average loss: 0.0002, Accuracy: 927/1000 (93%)

Train Epoch: 4	[0/6000 (0%)]	Loss: 0.217177
Train Epoch: 4	[640/6000 (11%)]	Loss: 0.099984
Train Epoch: 4	[1280/6000 (21%)]	Loss: 0.058514
Train Epoch: 4	[1920/6000 (32%)]	Loss: 0.125665
Train Epoch: 4	[2560/6000 (43%)]	Loss: 0.195694
Train Epoch: 4	[3200/6000 (53%)]	Loss: 0.113017
Train Epoch: 4	[3840/6000 (64%)]	Loss: 0.097260
Train Epoch: 4	[4480/6000 (74%)]	Loss: 0.123328
Train Epoch: 4	[5120/6000 (85%)]	Loss: 0.116845
Train Epoch: 4	[5760/6000 (96%)]	Loss: 0.093454

Test set: Average loss: 0.0001, Accuracy: 963/1000 (96%)

Train Epoch: 5	[0/6000 (0%)]	Loss: 0.094434
Train Epoch: 5	[640/6000 (11%)]	Loss: 0.196890
Train Epoch: 5	[1280/6000 (21%)]	Loss: 0.195391
Train Epoch: 5	[1920/6000 (32%)]	Loss: 0.162111
Train Epoch: 5	[2560/6000 (43%)]	Loss: 0.205471
Train Epoch: 5	[3200/6000 (53%)]	Loss: 0.321440
Train Epoch: 5	[3840/6000 (64%)]	Loss: 0.103642
Train Epoch: 5	[4480/6000 (74%)]	Loss: 0.091943
Train Epoch: 5	[5120/6000 (85%)]	Loss: 0.173878
Train Epoch: 5	[5760/6000 (96%)]	Loss: 0.075673

Test set: Average loss: 0.0001, Accuracy: 974/1000 (97%)

Train Epoch: 6	[0/6000 (0%)]	Loss: 0.078129
Train Epoch: 6	[640/6000 (11%)]	Loss: 0.133610
Train Epoch: 6	[1280/6000 (21%)]	Loss: 0.078885
Train Epoch: 6	[1920/6000 (32%)]	Loss: 0.119538
Train Epoch: 6	[2560/6000 (43%)]	Loss: 0.109704
Train Epoch: 6	[3200/6000 (53%)]	Loss: 0.117903



Train Epoch: 6 [3840/6000 (64%)]	Loss: 0.041985
Train Epoch: 6 [4480/6000 (74%)]	Loss: 0.055203
Train Epoch: 6 [5120/6000 (85%)]	Loss: 0.057204
Train Epoch: 6 [5760/6000 (96%)]	Loss: 0.041579

Test set: Average loss: 0.0001, Accuracy: 977/1000 (98%)

Train Epoch: 7 [0/6000 (0%)]	Loss: 0.064444
Train Epoch: 7 [640/6000 (11%)]	Loss: 0.059792
Train Epoch: 7 [1280/6000 (21%)]	Loss: 0.035186
Train Epoch: 7 [1920/6000 (32%)]	Loss: 0.135161
Train Epoch: 7 [2560/6000 (43%)]	Loss: 0.058195
Train Epoch: 7 [3200/6000 (53%)]	Loss: 0.129944
Train Epoch: 7 [3840/6000 (64%)]	Loss: 0.025679
Train Epoch: 7 [4480/6000 (74%)]	Loss: 0.213292
Train Epoch: 7 [5120/6000 (85%)]	Loss: 0.126770
Train Epoch: 7 [5760/6000 (96%)]	Loss: 0.164390

Test set: Average loss: 0.0000, Accuracy: 986/1000 (99%)

Train Epoch: 8 [0/6000 (0%)]	Loss: 0.027620
Train Epoch: 8 [640/6000 (11%)]	Loss: 0.050954
Train Epoch: 8 [1280/6000 (21%)]	Loss: 0.056885
Train Epoch: 8 [1920/6000 (32%)]	Loss: 0.053830
Train Epoch: 8 [2560/6000 (43%)]	Loss: 0.029372
Train Epoch: 8 [3200/6000 (53%)]	Loss: 0.078245
Train Epoch: 8 [3840/6000 (64%)]	Loss: 0.010862
Train Epoch: 8 [4480/6000 (74%)]	Loss: 0.042567
Train Epoch: 8 [5120/6000 (85%)]	Loss: 0.018613
Train Epoch: 8 [5760/6000 (96%)]	Loss: 0.078957

Test set: Average loss: 0.0000, Accuracy: 983/1000 (98%)

Train Epoch: 9 [0/6000 (0%)]	Loss: 0.133632
Train Epoch: 9 [640/6000 (11%)]	Loss: 0.012117
Train Epoch: 9 [1280/6000 (21%)]	Loss: 0.147912
Train Epoch: 9 [1920/6000 (32%)]	Loss: 0.057928
Train Epoch: 9 [2560/6000 (43%)]	Loss: 0.017033
Train Epoch: 9 [3200/6000 (53%)]	Loss: 0.027200
Train Epoch: 9 [3840/6000 (64%)]	Loss: 0.064777
Train Epoch: 9 [4480/6000 (74%)]	Loss: 0.071504
Train Epoch: 9 [5120/6000 (85%)]	Loss: 0.069630
Train Epoch: 9 [5760/6000 (96%)]	Loss: 0.007540

Test set: Average loss: 0.0001, Accuracy: 980/1000 (98%)

[ ]:

## HW2\_task4

March 14, 2023

Using the code from <https://github.com/drgripa1/resnet-cifar10>, ResNet model was fitted on CIFAR10 dataset with different optimizer, SGD, SGD with momentum 0.9 and Adam, also learning rates, 0.1, 0.3 and 0.5. The code below shows the instance of ResNet model with momentum SGD and learning rate of 0.5. I collected data of each case by replacing the “lr” value in `parse_args_train` function and optimizer used in class `ResNetModel`.

From the plots at the end, we can clearly see that learning rate of 0.1 performs best in momentum SGD and Adam, and learning rate of 0.1 performs best in SGD. Learning rate of 0.5 performs worst for all three optimizers as large learning rate leads to instability and overfitting, and may overshoot the optimal.

From the test train error recorded below, we can see that momentum sgd with learning rate of 0.1 has the lowest error, but Adam with learning rate of 0.5 has the highest test error. The reason behind is that ResNet is a deep neural network with many layers, and training it requires a significant amount of computational resources. Momentum SGD is computationally efficient because it only uses the first-order gradient information and has a simple update rule that requires fewer computations than Adam, which uses both first and second-order gradient information. CIFAR10 is also a relatively simple image classification dataset with well-defined classes and consistent patterns. Adam’s adaptive learning rate and momentum can help in situations where the gradient is noisy or the loss landscape is complex, but in the case of CIFAR10, momentum SGD may be sufficient for finding a good solution.

Test error			
	sgd	momentum sgd	Adam
0.1	41.45%	36.62%	49.04%
0.3	39.37%	48.51%	55.84%
0.5	40.22%	58.54%	60.60%

Training error			
	sgd	momentum sgd	Adam
0.1	44.12%	35.99%	48.58%
0.3	37.19%	54.54%	56.97%
0.5	42.68%	70.17%	56.21%

```
[1]: import random
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
import os
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import argparse
```

```
[2]: def parse_args_base(parser):
    parser.add_argument('--n', help='network depth', type=int, default=3)
    parser.add_argument('--batch', help='batch size', type=int, default=128)
    parser.add_argument('--dataset_dir', default='./dataset')
    return parser

def parse_args_train():
    parser = argparse.ArgumentParser()
    parser = parse_args_base(parser)
    parser.add_argument('--checkpoint_dir', default='./dataset')
    parser.add_argument('--print_freq', help='print loss freq', type=int,
↪ default=10)
    parser.add_argument('--save_params_freq', help='parameters saving freq',
↪ type=int, default=1000)
    parser.add_argument('--lr', help='initial learning rate', type=float,
↪ default=0.5) # change the initial learning rate to 0.3 and 0.5
    parser.add_argument('--momentum', help='optimizer momentum', type=float,
↪ default=0.9)
    parser.add_argument('--weight_decay', help='optimizer weight decay (L2 reg.
↪)', type=float, default=0.0001)
    parser.add_argument('--decay_lr_1', help='iteration at which lr decays 1st',
↪ type=int, default=400)
    parser.add_argument('--decay_lr_2', help='iteration at which lr decays 2nd',
↪ type=int, default=800)
    parser.add_argument('--lr_decay_rate', help='lr *= lr_decay_rate at
↪ decay_lr_i-th iteration', type=float, default=0.1)
    parser.add_argument('--n_iter', help='learning iterations', type=int,
↪ default=1000)
    parser.add_argument('-f')
    args = parser.parse_args()

    return args
```

```

def parse_args_test():
    parser = argparse.ArgumentParser()
    parser = parse_args_base(parser)
    parser.add_argument('--params_path', help='path to saved model weights',
↪default='./dataset/model_final.pth')
    parser.add_argument('-f')
    args = parser.parse_args()

    return args

```

```

[3]: def preprocess_train(tensor):
    tensor -= tensor.mean().item()
    tensor = F.pad(tensor, pad=(4, 4, 4, 4), mode='constant', value=0)
    t = random.randrange(8)
    l = random.randrange(8)
    tensor = tensor[:, t:t+32, l:l+32]
    if random.random() < 0.5:
        tensor = transforms.functional.hflip(tensor)

    return tensor

def preprocess_test(tensor):
    tensor -= tensor.mean().item()
    return tensor

def get_dataloader(is_train, batch_size, path):
    if is_train:
        return DataLoader(
            datasets.CIFAR10(path,
                             train=True,
                             download=True,
                             transform=transforms.Compose([
                                 transforms.ToTensor(),
                                 transforms.Lambda(preprocess_train)
                             ])),
            batch_size=batch_size,
            shuffle=True
        )
    else:
        return DataLoader(
            datasets.CIFAR10(path,
                             train=False,
                             download=True,
                             transform=transforms.Compose([
                                 transforms.ToTensor(),

```

```

        transforms.Lambda(preprocess_test)
    ])),
    batch_size=batch_size,
    shuffle=False
)

```

```

[4]: def init_weights(net, gain=0.02):
    def init_func(m):
        classname = m.__class__.__name__
        if hasattr(m, 'weight') and (classname.find('Conv') != -1 or classname.
↪find('Linear') != -1):
            nn.init.kaiming_normal_(m.weight.data, a=0, mode='fan_in')
            if hasattr(m, 'bias') and m.bias is not None:
                nn.init.constant_(m.bias.data, 0.0)
        elif classname.find('BatchNorm2d') != -1:
            nn.init.normal_(m.weight.data, 1.0, gain)
            nn.init.constant_(m.bias.data, 0.0)

    net.apply(init_func)

class ResNetModel:
    def __init__(self, opt, train=True):
        self.net = ResNetCifar(opt.n)
        if train:
            self.net.train()
        else:
            self.net.eval()
        if torch.cuda.is_available():
            self.device = torch.device('cuda')
            self.net = self.net.to('cuda')
            self.net = torch.nn.DataParallel(self.net)
        else:
            self.device = torch.device('cpu')

        init_weights(self.net)
        num_params = 0
        for param in self.net.parameters():
            num_params += param.numel()
        print(f'Total number of parameters : {num_params / 1e6:.3f} M')

        if train:
            self.checkpoint_dir = opt.checkpoint_dir
            self.optimizer = optim.SGD(          # replace with SGD , Adam
                self.net.parameters(),
                momentum=opt.momentum,
                lr=opt.lr,

```

```

        weight_decay=opt.weight_decay
    )
    self.scheduler = optim.lr_scheduler.MultiStepLR(
        self.optimizer,
        milestones=[opt.decay_lr_1, opt.decay_lr_2],
        gamma=opt.lr_decay_rate
    )
    self.criterion = nn.CrossEntropyLoss()
    self.loss = 0.0

def optimize_params(self, x, label):
    x = x.to(self.device)
    label = label.to(self.device)
    y = self._forward(x)
    self._update_params(y, label)

def _forward(self, x):
    return self.net(x)

def _backward(self, y, label):
    self.loss = self.criterion(y, label)
    self.loss.backward()

def _update_params(self, y, label):
    self.optimizer.zero_grad()
    self._backward(y, label)
    self.optimizer.step()
    self.scheduler.step() # scheduler step in each iteration

def test(self, x, label):
    with torch.no_grad():
        x = x.to(self.device)
        label = label.to(self.device)
        outputs = self._forward(x)
        _, predicted = torch.max(outputs.data, 1)
        total = label.size(0)
        correct = (predicted == label).sum().item()
        return correct, total, predicted

def val(self, x, label):
    with torch.no_grad():
        x = x.to(self.device)
        label = label.to(self.device)
        y = self._forward(x)
        return self.criterion(y, label).item()

def save_model(self, name):

```

```

path = os.path.join(self.checkpoint_dir, f'model_{name}.pth')
torch.save(self.net.state_dict(), path)
print(f'model saved to {path}')

def load_model(self, path):
    self.net.load_state_dict(torch.load(path))
    print(f'model loaded from {path}')

def get_current_loss(self):
    return self.loss.item()

```

```

[5]: class ResNetCifarBlock(nn.Module):
    def __init__(self, input_nc, output_nc):
        super().__init__()
        stride = 1
        self.expand = False
        if input_nc != output_nc:
            assert input_nc * 2 == output_nc, 'output_nc must be input_nc * 2'
            stride = 2
            self.expand = True

        self.conv1 = nn.Conv2d(input_nc, output_nc, kernel_size=3,
→stride=stride, padding=1)
        self.bn1 = nn.BatchNorm2d(output_nc)
        self.conv2 = nn.Conv2d(output_nc, output_nc, kernel_size=3, stride=1,
→padding=1)
        self.bn2 = nn.BatchNorm2d(output_nc)

    def forward(self, x):
        xx = F.relu(self.bn1(self.conv1(x)), inplace=True)
        y = self.bn2(self.conv2(xx))
        if self.expand:
            x = F.interpolate(x, scale_factor=0.5, mode='nearest') # subsampling
            zero = torch.zeros_like(x)
            x = torch.cat([x, zero], dim=1) # option A in the original paper
            h = F.relu(y + x, inplace=True)
            return h

def make_resblock_group(cls, input_nc, output_nc, n):
    blocks = []
    blocks.append(cls(input_nc, output_nc))
    for _ in range(1, n):
        blocks.append(cls(output_nc, output_nc))
    return nn.Sequential(*blocks)

```



```

class ResNetCifar(nn.Module):
    def __init__(self, n):
        super().__init__()

        self.conv = nn.Conv2d(3, 16, kernel_size=3, stride=1, padding=1)
        self.bn = nn.BatchNorm2d(16)
        self.block1 = make_resblock_group(ResNetCifarBlock, 16, 16, n)
        self.block2 = make_resblock_group(ResNetCifarBlock, 16, 32, n)
        self.block3 = make_resblock_group(ResNetCifarBlock, 32, 64, n)
        self.pool = nn.AdaptiveAvgPool2d(output_size=(1, 1)) # global average
    →pooling
        self.fc = nn.Linear(64, 10)

    def forward(self, x):
        x = F.relu(self.bn(self.conv(x)), inplace=True)
        x = self.block1(x)
        x = self.block2(x)
        x = self.block3(x)
        x = self.pool(x)
        x = x.view(x.shape[0], -1)
        x = self.fc(x)
        return x

```

```

[6]: def train(opt):
    dataloader = get_dataloader(True, opt.batch, opt.dataset_dir)
    model = ResNetModel(opt, train=True)
    print(opt)
    with open(os.path.join(opt.checkpoint_dir, 'loss_log.txt'), 'a') as f:
        f.write(str(opt) + '\n')

    total_iter = 0
    loss = 0.0

    while True:
        for batch in dataloader:
            total_iter += 1
            inputs, labels = batch
            model.optimize_params(inputs, labels)
            loss += model.get_current_loss()

            if total_iter % opt.print_freq == 0:
                txt = f'iter: {total_iter: 6d}, loss: {loss / opt.print_freq}'
                print(txt)
                txxt= str(loss / opt.print_freq)
                with open(os.path.join(opt.checkpoint_dir, 'loss_log.txt'), 'a')
    →as f:
                    f.write(txxt + '\n')

```

```

        loss = 0.0

    if total_iter % opt.save_params_freq == 0:
        model.save_model(f'{total_iter // opt.save_params_freq}k')

    if total_iter == opt.n_iter:
        model.save_model('final')
        return

if __name__ == '__main__':
    args = parse_args_train()
    train(args)

```

Files already downloaded and verified

Total number of parameters : 0.270 M

Namespace(batch=128, checkpoint\_dir='./dataset', dataset\_dir='./dataset',  
decay\_lr\_1=400, decay\_lr\_2=800, f='/root/.local/share/jupyter/runtime/kernel-  
48e8fd21-6e3d-43b5-a0d7-0557d02feaaa.json', lr=0.5, lr\_decay\_rate=0.1,  
momentum=0.9, n=3, n\_iter=1000, print\_freq=10, save\_params\_freq=1000,  
weight\_decay=0.0001)

```

iter:    10, loss: 5.260178208351135
iter:    20, loss: 2.386072850227356
iter:    30, loss: 2.317230987548828
iter:    40, loss: 2.3083648681640625
iter:    50, loss: 2.3303656816482543
iter:    60, loss: 2.314922499656677
iter:    70, loss: 2.3095060348510743
iter:    80, loss: 2.3009209394454957
iter:    90, loss: 2.3232123374938967
iter:   100, loss: 2.3156339168548583
iter:   110, loss: 2.3049930334091187
iter:   120, loss: 2.311630296707153
iter:   130, loss: 2.313529062271118
iter:   140, loss: 2.3088917255401613
iter:   150, loss: 2.320504069328308
iter:   160, loss: 2.3122891902923586
iter:   170, loss: 2.3100955963134764
iter:   180, loss: 2.3066614151000975
iter:   190, loss: 2.3113537549972536
iter:   200, loss: 2.312246012687683
iter:   210, loss: 2.309658408164978
iter:   220, loss: 2.303312659263611
iter:   230, loss: 2.301997923851013
iter:   240, loss: 2.2855573177337645
iter:   250, loss: 2.257064461708069
iter:   260, loss: 2.2389535188674925
iter:   270, loss: 2.2505528211593626

```

iter: 280, loss: 2.224766397476196  
iter: 290, loss: 2.1733654499053956  
iter: 300, loss: 2.173841452598572  
iter: 310, loss: 2.137011480331421  
iter: 320, loss: 2.157132863998413  
iter: 330, loss: 2.137747859954834  
iter: 340, loss: 2.09818320274353  
iter: 350, loss: 2.0846671581268312  
iter: 360, loss: 2.0936472058296203  
iter: 370, loss: 2.0883981227874755  
iter: 380, loss: 2.0420836567878724  
iter: 390, loss: 1.9952819108963014  
iter: 400, loss: 1.989047122001648  
iter: 410, loss: 1.9662629127502442  
iter: 420, loss: 1.9113463044166565  
iter: 430, loss: 1.9285452604293822  
iter: 440, loss: 1.891862964630127  
iter: 450, loss: 1.9047509908676148  
iter: 460, loss: 1.9072709918022155  
iter: 470, loss: 1.9212000250816346  
iter: 480, loss: 1.8909221172332764  
iter: 490, loss: 1.926684558391571  
iter: 500, loss: 1.9137438654899597  
iter: 510, loss: 1.9010217308998107  
iter: 520, loss: 1.908447551727295  
iter: 530, loss: 1.903776264190674  
iter: 540, loss: 1.9186978578567504  
iter: 550, loss: 1.880186128616333  
iter: 560, loss: 1.8685925364494325  
iter: 570, loss: 1.8971842050552368  
iter: 580, loss: 1.8854424595832824  
iter: 590, loss: 1.8786907196044922  
iter: 600, loss: 1.8689462304115296  
iter: 610, loss: 1.8682750463485718  
iter: 620, loss: 1.8887632846832276  
iter: 630, loss: 1.8512694716453553  
iter: 640, loss: 1.8586389422416687  
iter: 650, loss: 1.8563894271850585  
iter: 660, loss: 1.8380138993263244  
iter: 670, loss: 1.8822349309921265  
iter: 680, loss: 1.8472273111343385  
iter: 690, loss: 1.8295931100845337  
iter: 700, loss: 1.8575900197029114  
iter: 710, loss: 1.861980676651001  
iter: 720, loss: 1.8861384987831116  
iter: 730, loss: 1.83902450799942  
iter: 740, loss: 1.8680475950241089  
iter: 750, loss: 1.8497129201889038

```

iter: 760, loss: 1.8317009091377259
iter: 770, loss: 1.8459845781326294
iter: 780, loss: 1.8697036504745483
iter: 790, loss: 1.8705100774765016
iter: 800, loss: 1.8529508590698243
iter: 810, loss: 1.8222689867019652
iter: 820, loss: 1.8134443640708924
iter: 830, loss: 1.8534798622131348
iter: 840, loss: 1.8102815628051758
iter: 850, loss: 1.8319310426712037
iter: 860, loss: 1.8184168219566346
iter: 870, loss: 1.8312573194503785
iter: 880, loss: 1.7968510746955872
iter: 890, loss: 1.8029914617538452
iter: 900, loss: 1.8617348432540894
iter: 910, loss: 1.7939058899879456
iter: 920, loss: 1.8065063953399658
iter: 930, loss: 1.8441028594970703
iter: 940, loss: 1.8144920110702514
iter: 950, loss: 1.8326905369758606
iter: 960, loss: 1.8100410938262939
iter: 970, loss: 1.8593904733657838
iter: 980, loss: 1.8339136362075805
iter: 990, loss: 1.816961658000946
iter: 1000, loss: 1.8099796295166015
model saved to ./dataset/model_1k.pth
model saved to ./dataset/model_final.pth

```

```

[7]: def trainerror(opt):
    print(opt)
    dataloader = get_dataloader(True, opt.batch, opt.dataset_dir)
    model = ResNetModel(opt, train=False)
    model.load_model(opt.params_path)

    total_n = 0
    total_correct = 0

    for batch in dataloader:
        inputs, labels = batch
        correct, total, _ = model.test(inputs, labels)
        total_correct += correct
        total_n += total

    acc = 100 * total_correct / total_n
    err = 100 - acc
    print(f'accuracy: {acc:.2f} %')
    print(f'error: {err:.2f} %')

```

```

print(f'{total_correct} / {total_n}')

if __name__ == '__main__':
    args = parse_args_test()
    trainerror(args)

```

```

Namespace(batch=128, dataset_dir='./dataset', f='/root/.local/share/jupyter/runtime/kernel-48e8fd21-6e3d-43b5-a0d7-0557d02feaaa.json', n=3,
params_path='./dataset/model_final.pth')
Files already downloaded and verified
Total number of parameters : 0.270 M
model loaded from ./dataset/model_final.pth
accuracy: 29.83 %
error: 70.17 %
14913 / 50000

```

```

[8]: def test(opt):
    dataloader = get_dataloader(False, opt.batch, opt.dataset_dir)
    model = ResNetModel(opt, train=True)
    print(opt)
    with open(os.path.join(opt.checkpoint_dir, 'test_loss_log.txt'), 'a') as f:
        f.write(str(opt) + '\n')

    total_iter = 0
    loss = 0.0

    while True:
        for batch in dataloader:
            total_iter += 1
            inputs, labels = batch
            model.optimize_params(inputs, labels)
            loss += model.get_current_loss()

            if total_iter % opt.print_freq == 0:
                txt = f'iter: {total_iter: 6d}, loss: {loss / opt.print_freq}'
                print(txt)
                txxt= str(loss / opt.print_freq)
                with open(os.path.join(opt.checkpoint_dir, 'test_loss_log.txt'),
→ 'a') as f:
                    f.write(txxt + '\n')

                loss = 0.0

            if total_iter % opt.save_params_freq == 0:
                model.save_model(f'{total_iter // opt.save_params_freq}k')

```

```

        if total_iter == opt.n_iter:
            model.save_model('final')
            return

if __name__ == '__main__':
    args = parse_args_train()
    train(args)

```

Files already downloaded and verified

Total number of parameters : 0.270 M

Namespace(batch=128, checkpoint\_dir='./dataset', dataset\_dir='./dataset',  
decay\_lr\_1=400, decay\_lr\_2=800, f='/root/.local/share/jupyter/runtime/kernel-  
48e8fd21-6e3d-43b5-a0d7-0557d02feaaa.json', lr=0.5, lr\_decay\_rate=0.1,  
momentum=0.9, n=3, n\_iter=1000, print\_freq=10, save\_params\_freq=1000,  
weight\_decay=0.0001)

```

iter:    10, loss: 5.180495095252991
iter:    20, loss: 2.5708837032318117
iter:    30, loss: 2.3134934663772584
iter:    40, loss: 2.294820785522461
iter:    50, loss: 2.305798149108887
iter:    60, loss: 2.235015892982483
iter:    70, loss: 2.202967810630798
iter:    80, loss: 2.1293555736541747
iter:    90, loss: 2.0789914965629577
iter:   100, loss: 2.070948827266693
iter:   110, loss: 2.0869250297546387
iter:   120, loss: 2.021182823181152
iter:   130, loss: 2.0224040389060973
iter:   140, loss: 2.033132183551788
iter:   150, loss: 2.012133979797363
iter:   160, loss: 1.9620522499084472
iter:   170, loss: 1.9950168967247008
iter:   180, loss: 1.9844149231910706
iter:   190, loss: 1.9728642225265502
iter:   200, loss: 1.9004098892211914
iter:   210, loss: 1.9750344157218933
iter:   220, loss: 1.9065344333648682
iter:   230, loss: 1.9135539293289185
iter:   240, loss: 1.9360596179962157
iter:   250, loss: 1.8860053777694703
iter:   260, loss: 1.8376816630363464
iter:   270, loss: 1.9220985889434814
iter:   280, loss: 1.8785597801208496
iter:   290, loss: 1.8995963215827942
iter:   300, loss: 1.8678790211677552

```

iter: 310, loss: 1.8353885531425476  
iter: 320, loss: 1.8193363785743712  
iter: 330, loss: 1.7869680166244506  
iter: 340, loss: 1.8425472140312196  
iter: 350, loss: 1.8245766162872314  
iter: 360, loss: 1.8276852250099183  
iter: 370, loss: 1.7972840666770935  
iter: 380, loss: 1.87012220621109  
iter: 390, loss: 1.7960012316703797  
iter: 400, loss: 1.7840522289276124  
iter: 410, loss: 1.8198033213615417  
iter: 420, loss: 1.7096255421638489  
iter: 430, loss: 1.6872313261032104  
iter: 440, loss: 1.7210400342941283  
iter: 450, loss: 1.6903883457183837  
iter: 460, loss: 1.7191497921943664  
iter: 470, loss: 1.705140483379364  
iter: 480, loss: 1.6766711950302124  
iter: 490, loss: 1.659067404270172  
iter: 500, loss: 1.7093275427818297  
iter: 510, loss: 1.6726937770843506  
iter: 520, loss: 1.6716200590133667  
iter: 530, loss: 1.6829920291900635  
iter: 540, loss: 1.6174623489379882  
iter: 550, loss: 1.676980459690094  
iter: 560, loss: 1.7139034628868104  
iter: 570, loss: 1.6504444837570191  
iter: 580, loss: 1.5975868344306945  
iter: 590, loss: 1.6268171072006226  
iter: 600, loss: 1.681311798095703  
iter: 610, loss: 1.6315538048744203  
iter: 620, loss: 1.6422980546951294  
iter: 630, loss: 1.621252679824829  
iter: 640, loss: 1.6520115852355957  
iter: 650, loss: 1.6609533548355102  
iter: 660, loss: 1.647756004333496  
iter: 670, loss: 1.6445759057998657  
iter: 680, loss: 1.6381321549415588  
iter: 690, loss: 1.5924819350242614  
iter: 700, loss: 1.6437740325927734  
iter: 710, loss: 1.59366112947464  
iter: 720, loss: 1.6274990439414978  
iter: 730, loss: 1.6081182718276978  
iter: 740, loss: 1.6314485669136047  
iter: 750, loss: 1.5685015320777893  
iter: 760, loss: 1.6424452424049378  
iter: 770, loss: 1.6488170385360719  
iter: 780, loss: 1.641170573234558

```

iter:    790, loss: 1.5890581130981445
iter:    800, loss: 1.6045084834098815
iter:    810, loss: 1.6279166221618653
iter:    820, loss: 1.6227516770362853
iter:    830, loss: 1.6034451007843018
iter:    840, loss: 1.590772521495819
iter:    850, loss: 1.6050360321998596
iter:    860, loss: 1.6205818057060242
iter:    870, loss: 1.615891981124878
iter:    880, loss: 1.5819509744644165
iter:    890, loss: 1.5619788646697998
iter:    900, loss: 1.5797269225120545
iter:    910, loss: 1.6152411460876466
iter:    920, loss: 1.596661376953125
iter:    930, loss: 1.582091212272644
iter:    940, loss: 1.5766264200210571
iter:    950, loss: 1.557748556137085
iter:    960, loss: 1.5945794343948365
iter:    970, loss: 1.5818885684013366
iter:    980, loss: 1.5264253616333008
iter:    990, loss: 1.5639222264289856
iter:   1000, loss: 1.5881321787834168
model saved to ./dataset/model_1k.pth
model saved to ./dataset/model_final.pth

```

```

[9]: def testerror(opt):
    print(opt)
    dataloader = get_dataloader(False, opt.batch, opt.dataset_dir)
    model = ResNetModel(opt, train=False)
    model.load_model(opt.params_path)

    total_n = 0
    total_correct = 0

    for batch in dataloader:
        inputs, labels = batch
        correct, total, _ = model.test(inputs, labels)
        total_correct += correct
        total_n += total

    acc = 100 * total_correct / total_n
    err = 100 - acc
    print(f'accuracy: {acc:.2f} %')
    print(f'error: {err:.2f} %')
    print(f'{total_correct} / {total_n}')

```



```

if __name__ == '__main__':
    args = parse_args_test()
    testerror(args)

```

```

Namespace(batch=128, dataset_dir='./dataset', f='/root/.local/share/jupyter/runtime/kernel-48e8fd21-6e3d-43b5-a0d7-0557d02feaaa.json', n=3,
params_path='./dataset/model_final.pth')
Files already downloaded and verified
Total number of parameters : 0.270 M
model loaded from ./dataset/model_final.pth
accuracy: 41.46 %
error: 58.54 %
4146 / 10000

```

```

[10]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
split=100
df=pd.read_csv("./dataset/loss_log.txt", skiprows=[0,split+1], header=None)
np.savetxt(r'./dataset/msgd_5_train.txt', df[:split].values)
np.savetxt(r'./dataset/msgd_5_test.txt', df[split:].values)

```

```

[13]: msgd_5_train=pd.read_csv("./dataset/msgd_5_train.txt", header=None)
msgd_5_test=pd.read_csv("./dataset/msgd_5_test.txt", header=None)
msgd_3_train=pd.read_csv("./dataset/msgd_3_train.txt", header=None)
msgd_3_test=pd.read_csv("./dataset/msgd_3_test.txt", header=None)
msgd_1_train=pd.read_csv("./dataset/msgd_1_train.txt", header=None)
msgd_1_test=pd.read_csv("./dataset/msgd_1_test.txt", header=None)

sgd_5_train=pd.read_csv("./dataset/sgd_5_train.txt", header=None)
sgd_5_test=pd.read_csv("./dataset/sgd_5_test.txt", header=None)
sgd_3_train=pd.read_csv("./dataset/sgd_3_train.txt", header=None)
sgd_3_test=pd.read_csv("./dataset/sgd_3_test.txt", header=None)
sgd_1_train=pd.read_csv("./dataset/sgd_1_train.txt", header=None)
sgd_1_test=pd.read_csv("./dataset/sgd_1_test.txt", header=None)

adam_5_test=pd.read_csv("./dataset/adam_5_test.txt", header=None)
adam_5_train=pd.read_csv("./dataset/adam_5_train.txt", header=None)
adam_3_test=pd.read_csv("./dataset/adam_3_test.txt", header=None)
adam_3_train=pd.read_csv("./dataset/adam_3_train.txt", header=None)
adam_1_test=pd.read_csv("./dataset/adam_1_test.txt", header=None)
adam_1_train=pd.read_csv("./dataset/adam_1_train.txt", header=None)


```

```


[14]: x=np.linspace(0,1000,split)
plt.plot(x, msgd_5_train, label='training loss with 0.5 learning rate')
plt.plot(x, msgd_5_test, label='test loss with 0.5 learning rate')
plt.plot(x, msgd_3_train, label='training loss with 0.3 learning rate')
plt.plot(x, msgd_3_test, label='test loss with 0.3 learning rate')

```

```
plt.plot(x, msgd_1_train, label='training loss with 0.1 learning rate')
plt.plot(x, msgd_1_test, label='test loss with 0.1 learning rate')
plt.ylabel('Loss')
plt.xlabel('Iterations')
plt.legend(loc='upper right')
plt.title("ResNet with momentum SGD optimizer")
plt.show()
```



```
[15]: plt.plot(x, sgd_5_train, label='training loss with 0.5 learning rate')
plt.plot(x, sgd_5_test, label='test loss with 0.5 learning rate')
plt.plot(x, sgd_3_train, label='training loss with 0.3 learning rate')
plt.plot(x, sgd_3_test, label='test loss with 0.3 learning rate')
plt.plot(x, sgd_1_train, label='training loss with 0.1 learning rate')
plt.plot(x, sgd_1_test, label='test loss with 0.1 learning rate')
plt.ylabel('Loss')
plt.xlabel('Iterations')
plt.legend(loc='upper right')
plt.title("ResNet with SGD optimizer")
plt.show()
```



```
[16]: plt.plot(x, adam_5_train, label='training loss with 0.5 learning rate')
plt.plot(x, adam_5_test, label='test loss with 0.5 learning rate')
plt.plot(x, adam_3_train, label='training loss with 0.3 learning rate')
plt.plot(x, adam_3_test, label='test loss with 0.3 learning rate')
plt.plot(x, adam_1_train, label='training loss with 0.1 learning rate')
plt.plot(x, adam_1_test, label='test loss with 0.1 learning rate')
plt.ylabel('Loss')
plt.xlabel('Iterations')
plt.legend(loc='upper right')
plt.title("ResNet with Adam optimizer")
plt.show()
```

